

Programming Assignment #1

Multi-Core MapReduce

CS 416: Operating System Design
Due 11:55 PM, October 6

Professor: Ricardo Bianchini
TA: William Katsak

Group Size: 3
Programming Language: **C or C++ ONLY**

Overview

For this assignment, you will implement a framework for multithreaded data processing on a single computer that follows the "MapReduce" paradigm.

MapReduce is a strategy for processing parallelizable problems across large datasets using a large number of processing elements (either multiple CPU cores or multiple individual computers).

The name MapReduce comes from the two-phase manner in which the data is processed:

In the **map** phase, the framework takes the input, partitions it into smaller subproblems (subsets of the input), and distributes these subproblems to individual workers. Each worker processes its assigned subproblem, and returns its output to the framework.

In the **reduce** phase, the framework collects all of the **map** outputs (i.e. the map phase must have finished) and combines them in some way to form the output to the problem that is being solved.

Assignment

1. You are required to design and implement a MapReduce framework using C or C++ on Linux.
2. Your framework must be configurable to provide parallelism in two distinct ways:
 - a. Using multiple Linux processes
 - b. Using POSIX threads (pthreads)
3. Your framework should be designed in a general manner. In particular, the problem to be solved (see item 5 below) should be encapsulated solely in the implementation of the `map()` and `reduce()` functions. Substitution of different `map()` and `reduce()` functions should enable your framework to solve a different problem without ***any further changes in the framework itself***.
4. All intermediate data that is to be transferred between the map and reduce phases is to be stored in **POSIX shared memory**. Reading and writing to this shared memory may require synchronization (mutual exclusion). You may NOT write intermediate output to the file system. There also needs to be synchronization (barrier) to explicitly separate the map and reduce phases.
5. Once your framework is working, you must implement two sets of map/reduce functions to solve the following problems:
 - a. **Word count**: given an input file, you must count the number of times each individual word appears in the input, and output a file containing a list of words followed by their counts.
 - b. **Integer sort**: given an input file containing a list of integers, you must output a file containing the same integers in sorted, ascending order.
6. You will be provided with sample input and output files for each problem. Your implementation must be able to process the example input and must produce output in **exactly** the same format as the example output file. (These sample files will be attached to the assignment on Sakai).

7. Your framework must compile to a single executable file called **mapred**, that conforms to the following command line structure:

```
mapred --app [wordcount, sort] --impl [procs, threads]
--maps num_maps --reduces num_reduces --input infile
--output outfile
```

e.g. You can execute this four different ways:

- wordcount with processes
- wordcount with threads
- sort with processes
- sort with threads

with any arbitrary number of maps and reduces.

Honors/Extra Credit

All students in the honors section must meet these additional requirements for full credit; students in the regular sections may earn up to 20% extra credit by meeting these requirements ***provided that they also meet all of the base requirements.***

1. You must implement an additional multithreaded solution for each of the specified problems (integer sort and wordcount) using a non-MapReduce parallel implementation.
2. Your additional implementation must use pthreads.
3. The additional solution must be activated by adding an additional option to the --impl flag called "extra". In this case, the --maps and --reduces arguments will not be necessary, and should be replaced by a single --numthreads [thread_count] argument.

For Example:

```
mapred --app wordcount --impl extra --numthreads 8
--input infile --output outfile
```

Report

You must also provide a report that contains, at the minimum:

1. The names of the group members.
2. An architectural description of your framework.
3. An evaluation of the performance difference between the process-based and thread-based implementations of your framework.
4. A description of any difficulties that you encountered during this project.

The report must be in PDF format to receive credit, **no exceptions**.

Submission

Solutions must be submitted to Sakai, attached to the assignment. Each group should nominate one member to submit their solution. Each student SHOULD NOT submit an individual copy of the solution.

Each submission must consist of two files:

1. A .tar file containing the source code and Makefile.
2. A PDF file containing the report (this should not be inside the .tar file).

Reference

- MapReduce:
 - <http://en.wikipedia.org/wiki/MapReduce>
- Mutual Exclusion and Synchronization
 - http://en.wikipedia.org/wiki/Mutual_exclusion
 - [http://en.wikipedia.org/wiki/Semaphore_\(programming\)](http://en.wikipedia.org/wiki/Semaphore_(programming))
- Creating and using multiple processes
 - <http://www.yolinux.com/TUTORIALS/ForkExecProcesses.html>
- Creating and using multiple threads
 - <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.htm>
- Using POSIX shared memory
 - <http://www.cs.cf.ac.uk/Dave/C/node27.html>
- Textbook
 - Operating System Concepts, Silberschatz and Galvin